

CensusPay.R

rharidas

2021-12-02

```
# Note: This script will take a while to run. In particular the knn and random forest algorithms with tuning grids will take  
# more time. please be patient if you happen to execute it. The execution report is available in the github location as well
```

```
# Execute the given source code for the project  
source("DatasetProcessingCode.R")
```

```
## Loading required package: tidyverse
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr   0.3.4  
## v tibble  3.1.2      v dplyr  1.0.7  
## v tidyr   1.1.3      v stringr 1.4.0  
## v readr   1.4.0      v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()    masks stats::lag()
```

```
## Loading required package: caret
```

```
## Loading required package: lattice
```

```
##  
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':  
##  
## lift
```

```
## Loading required package: data.table
```

```
##  
## Attaching package: 'data.table'
```

```
## The following objects are masked from 'package:dplyr':  
##  
## between, first, last
```

```
## The following object is masked from 'package:purrr':  
##  
## transpose
```

```
## Loading required package: gridExtra
```

```
##  
## Attaching package: 'gridExtra'
```

```
## The following object is masked from 'package:dplyr':  
##  
## combine
```

```
## Loading required package: kableExtra
```

```
##  
## Attaching package: 'kableExtra'
```

```
## The following object is masked from 'package:dplyr':  
##  
##   group_rows
```

```
## Loading required package: epiDisplay
```

```
## Loading required package: foreign
```

```
## Loading required package: survival
```

```
##  
## Attaching package: 'survival'
```

```
## The following object is masked from 'package:caret':  
##  
##   cluster
```

```
## Loading required package: MASS
```

```
##  
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':  
##  
##   select
```

```
## Loading required package: nnet
```

```
##  
## Attaching package: 'epiDisplay'
```

```
## The following object is masked from 'package:lattice':  
##  
##   dotplot
```

```
## The following object is masked from 'package:ggplot2':  
##  
##   alpha
```


```
## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler  
## used
```

```


## Rows: 32,561
## Columns: 15
## $ age          <int> 90, 82, 66, 54, 41, 34, 38, 74, 68, 41, 45, 38, 52, 32, ~
## $ workclass    <chr> "?", "Private", "?", "Private", "Private", "Private", "~
## $ fnlwt        <int> 77053, 132870, 186061, 140359, 264663, 216864, 150601, ~
## $ education    <chr> "HS-grad", "HS-grad", "Some-college", "7th-8th", "Some-~
## $ education.num <int> 9, 9, 10, 4, 10, 9, 6, 16, 9, 10, 16, 15, 13, 14, 16, 1~
## $ marital.status <chr> "Widowed", "Widowed", "Widowed", "Divorced", "Separated~
## $ occupation   <chr> "?", "Exec-managerial", "?", "Machine-op-inspct", "Prof~
## $ relationship <chr> "Not-in-family", "Not-in-family", "Unmarried", "Unmarri~
## $ race          <chr> "White", "White", "Black", "White", "White", "White", "~
## $ sex           <chr> "Female", "Female", "Female", "Female", "Female", "Fema~
## $ capital.gain  <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ capital.loss  <int> 4356, 4356, 4356, 3900, 3900, 3770, 3770, 3683, 3683, 3~
## $ hours.per.week <int> 40, 18, 40, 40, 40, 45, 40, 20, 40, 60, 35, 45, 20, 55, ~
## $ native.country <chr> "United-States", "United-States", "United-States", "Uni~
## $ income        <chr> "<=50K", "<=50K", "<=50K", "<=50K", "<=50K", "<=50K", "~
## Rows: 29,170
## Columns: 13
## $ age          <int> 90, 82, 66, 54, 41, 34, 38, 74, 68, 45, 38, 52, 32, 51, ~
## $ fnlwt        <int> 77053, 132870, 186061, 140359, 264663, 216864, 150601, 8~
## $ education    <fct> HSgrad, HSgrad, Somecollege, 7th8th, Somecollege, HSgrad~
## $ eduyears     <int> 9, 9, 10, 4, 10, 9, 6, 16, 9, 16, 15, 13, 14, 16, 15, 7, ~
## $ maritalstatus <fct> Widowed, Widowed, Widowed, Divorced, Separated, Divorced~
## $ occupation   <fct> Prof-specialty, Execmanagerial, Prof-specialty, Machineo~
## $ relationship <fct> Notinfamily, Notinfamily, Unmarried, Unmarried, Ownchild~
## $ race          <fct> White, White, Black, White, White, White, White, White, ~
## $ sex           <fct> Female, Female, Female, Female, Female, Female, Male, Fe~
## $ hoursperweek  <int> 40, 18, 40, 40, 40, 45, 40, 20, 40, 35, 45, 20, 55, 40, ~
## $ native       <chr> "UnitedStates", "UnitedStates", "UnitedStates", "UnitedS~
## $ income        <fct> AtBelow50K, AtBelow50K, AtBelow50K, AtBelow50K, AtBelow5~
## $ class         <fct> Private, Private, Private, Private, Private, Private, Pr~


```


Distribution of adultpayclean\$education

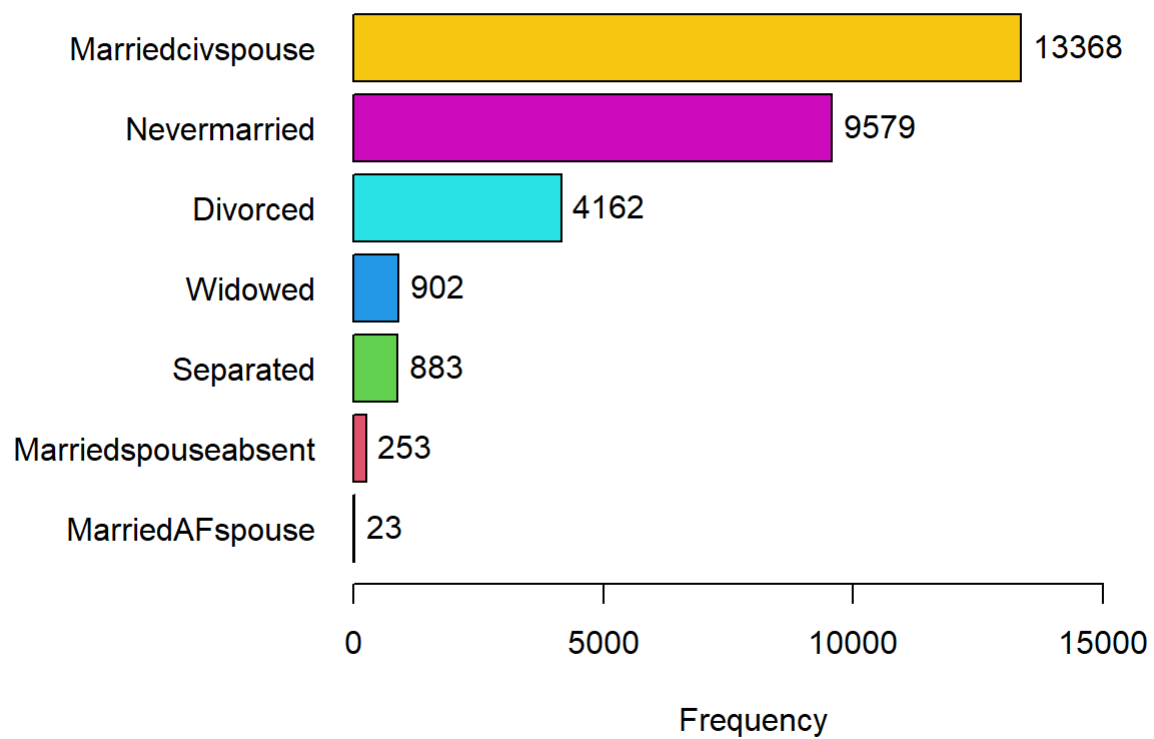


Distribution of adultpayclean\$race






Distribution of adultpayclean\$maritalstatus




Distribution of adultpayclean\$sex




Distribution of adultpayclean\$relationship



Distribution of adultpayclean\$class



Distribution of adultpayclean\$income



```
if (!require(randomForest))  
  install.packages("randomForest", repos = "http://cran.us.r-project.org")
```

```
## Loading required package: randomForest
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:gridExtra':  
##  
##      combine
```

```
## The following object is masked from 'package:dplyr':  
##  
##      combine
```

```
## The following object is masked from 'package:ggplot2':  
##  
##      margin
```

```
if (!require(purrr))  
  install.packages("purrr", repos = "http://cran.us.r-project.org")  
  
if (!require(e1071))  
  install.packages("e1071")
```

```
## Loading required package: e1071
```

```
if (!require(pROC))  
  install.packages("pRoc")
```

```
## Loading required package: pROC
```

```
## Type 'citation("pROC")' for a citation.
```

```
##  
## Attaching package: 'pROC'
```

```
## The following object is masked from 'package:epiDisplay':  
##  
##      ci
```

```
## The following objects are masked from 'package:stats':  
##  
##      cov, smooth, var
```

```
if (!require(ROCit))  
  install.packages("ROCit")
```

```
## Loading required package: ROCit
```

```
## Warning: package 'ROCit' was built under R version 4.1.2
```

```
library(caret)  
library(gridExtra)  
library(kableExtra)  
library(randomForest)  
library(purrr)  
library(e1071)  
library(caTools)  
library(pROC)  
library(ROCit)
```

```
#set the seed for reproducible results  
set.seed(2008, sample.kind = "Rounding")
```

```
## Warning in set.seed(2008, sample.kind = "Rounding"): non-uniform 'Rounding'  
## sampler used
```

```
# the simplest possible machine algorithm: guessing the outcome
seat_of_the_pants <-
  sample(c("Above50K", "AtBelow50K"), length(test_index), replace = TRUE) %>% factor(levels = levels(adultpayclean_validation$income))
# calculate the accuracy of this sampling
accuracy_guess <-
  mean(seat_of_the_pants == adultpayclean_validation$income)

# build a confusion matrix for this simple model
table(predicted = seat_of_the_pants, actual = adultpayclean_validation$income)
```

```
##           actual
## predicted  Above50K AtBelow50K
## Above50K      361      1078
## AtBelow50K     357      1122
```

```
# tabulate accuracy by income levels
adultpayclean_validation %>%
  mutate(y_hat = seat_of_the_pants) %>%
  group_by(income) %>%
  summarize(accuracy = mean(y_hat == income))
```

```
## # A tibble: 2 x 2
##   income      accuracy
##   <fct>      <dbl>
## 1 Above50K    0.503
## 2 AtBelow50K  0.51
```

```
# confusion matrix using R function
cm <-
  confusionMatrix(data = seat_of_the_pants , reference = adultpayclean_validation$income)
# display the confusion matrix
cm
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  Above50K AtBelow50K
##   Above50K      361      1078
##   AtBelow50K     357      1122
##
##           Accuracy : 0.5082
##           95% CI : (0.4899, 0.5265)
##   No Information Rate : 0.7539
##   P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0096
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.5028
##           Specificity : 0.5100
##   Pos Pred Value : 0.2509
##   Neg Pred Value : 0.7586
##   Prevalence : 0.2461
##   Detection Rate : 0.1237
##   Detection Prevalence : 0.4931
##   Balanced Accuracy : 0.5064
##
##   'Positive' Class : Above50K
##

```

```

#record the sensitivity, specificity, and prevalence
sensitivity_guess <- cm$byClass[["Sensitivity"]]
specificity_guess <- cm$byClass[["Specificity"]]
prevalence_guess <- cm$byClass[["Prevalence"]]
f1_guess <- cm$byClass[["F1"]]

```

```

#find the area under the curve/ROC
auc(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(seat_of_the_pants == "Above50K",1,0))

```



```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Area under the curve: 0.5064
```

```
set.seed(2008)
#logistic linear model
# create the model
lm_fit <- adu1tpayclean_train %>%
  mutate(y = as.numeric(income == "Above50K")) %>%
  lm(y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship + education,
     data = .)

# predict using test set
p_hat_logit <- predict(lm_fit, newdata = adu1tpayclean_validation)
```

```
## Warning in predict.lm(lm_fit, newdata = adu1tpayclean_validation): prediction
## from a rank-deficient fit may be misleading
```

```
#translate predicted data into factor
y_hat_logit <-
  ifelse(p_hat_logit > 0.5, "Above50K", "AtBelow50K") %>% factor

#compare the predicted vs observed values and use confusionMatrix to get the accuracy and other metrics
cm_lm <-
  confusionMatrix(y_hat_logit, adu1tpayclean_validation$income)
accuracy_lm <-
  confusionMatrix(y_hat_logit, adu1tpayclean_validation$income)$overall[["Accuracy"]]

cm_lm
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  Above50K AtBelow50K
##   Above50K      318      139
##   AtBelow50K     400     2061
##
##           Accuracy : 0.8153
##           95% CI : (0.8007, 0.8292)
##   No Information Rate : 0.7539
##   P-Value [Acc > NIR] : 1.243e-15
##
##           Kappa : 0.4327
##
##   Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.4429
##           Specificity : 0.9368
##   Pos Pred Value : 0.6958
##   Neg Pred Value : 0.8375
##   Prevalence : 0.2461
##   Detection Rate : 0.1090
##   Detection Prevalence : 0.1566
##   Balanced Accuracy : 0.6899
##
##   'Positive' Class : Above50K
##
```

```
#record the sensitivity, specificity, and prevalence
```

```
sensitivity_lm <- cm_lm$byClass[["Sensitivity"]]
```

```
specificity_lm <- cm_lm$byClass[["Specificity"]]
```

```
prevalence_lm <- cm_lm$byClass[["Prevalence"]]
```

```
f1_lm <- cm_lm$byClass[["F1"]]
```


```
#Find the ROC and plot it. Show the AUC as well
```

```
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_logit) == "Above50K",1,0),method="bin")
```

```
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)
```

```
plot(ciROC_bin95, col = 1, values=TRUE)
```

```
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)
```



```
ROCit::ciAUC(pROC_bin)
```

```
##  
## estimated AUC : 0.81709434414188  
## AUC estimation method : binormal  
##  
## CI of AUC  
## confidence level = 95%  
## lower = 0.792575444461593 upper = 0.841613243822167
```

```
set.seed(2008)  
#general linear model  
#create the glm model  
glm_fit <- adultpayclean_train %>%  
  mutate(y = as.numeric(income == "Above50K")) %>%  
  glm(  
    y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship + education,  
    data = .,  
    family = "binomial"  
  )  
  
# predict using validation set  
p_hat_logit <- predict(glm_fit, newdata = adultpayclean_validation)
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :  
## prediction from a rank-deficient fit may be misleading
```

```
# translate the predicted data into factor
y_hat_logit <-
  ifelse(p_hat_logit > 0.5, "Above50K", "AtBelow50K") %>% factor

# compare the predicted vs observed values and use confusionMatrix to get the accuracy and other metrics for the glm model
cm_glm <-
  confusionMatrix(y_hat_logit, adultpayclean_validation$income)
accuracy_glm <-
  confusionMatrix(y_hat_logit, adultpayclean_validation$income)$overall[["Accuracy"]]

cm_glm
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  Above50K AtBelow50K
##   Above50K      278      113
##   AtBelow50K    440     2087
##
##           Accuracy : 0.8105
##           95% CI : (0.7958, 0.8246)
##   No Information Rate : 0.7539
##   P-Value [Acc > NIR] : 1.758e-13
##
##           Kappa : 0.3967
##
##   McNemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.38719
##           Specificity : 0.94864
##   Pos Pred Value : 0.71100
##   Neg Pred Value : 0.82588
##   Prevalence : 0.24606
##   Detection Rate : 0.09527
##   Detection Prevalence : 0.13400
##   Balanced Accuracy : 0.66791
##
##   'Positive' Class : Above50K
##
```

```

#record the sensitivity, specificity, and prevalence
sensitivity_glm <- cm_glm$byClass[["Sensitivity"]]
specificity_glm <- cm_glm$byClass[["Specificity"]]
prevalence_glm <- cm_glm$byClass[["Prevalence"]]
f1_glm <- cm_glm$byClass[["F1"]]

#Find the ROC and plot it. Show the AUC as well
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_logit) == "Above50K",
,1,0),method="bin")
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)
plot(ciROC_bin95, col = 1, values=TRUE)
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)

```



```
ROCit::ciAUC(pROC_bin)
```

```
##  
## estimated AUC : 0.818172245714073  
## AUC estimation method : binormal  
##  
## CI of AUC  
## confidence level = 95%  
## lower = 0.791852436696456 upper = 0.844492054731691
```

```
#Naive bayes  
set.seed(2008)  
#create the naive bayes model  
train_nb <- adu1tpayclean_train %>%  
  mutate(y = as.factor(income == "Above50K")) %>%  
  naiveBayes(y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship+education,data = .)  
  
#predict using the validation dataset  
y_hat_nb <- predict(train_nb, newdata = adu1tpayclean_validation)  
#create the confusion matrix  
cm_tab <- table(adu1tpayclean_validation$income == "Above50K", y_hat_nb)  
cm_nb <- confusionMatrix(cm_tab)  
cm_nb
```




```
## Confusion Matrix and Statistics
##
##      y_hat_nb
##      FALSE TRUE
## FALSE 1794 406
## TRUE   175 543
##
##      Accuracy : 0.8009
##      95% CI : (0.7859, 0.8152)
##      No Information Rate : 0.6748
##      P-Value [Acc > NIR] : < 2.2e-16
##
##      Kappa : 0.5158
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 0.9111
##      Specificity : 0.5722
##      Pos Pred Value : 0.8155
##      Neg Pred Value : 0.7563
##      Prevalence : 0.6748
##      Detection Rate : 0.6148
##      Detection Prevalence : 0.7539
##      Balanced Accuracy : 0.7417
##
##      'Positive' Class : FALSE
##
```

#get the accuracy, sensitivity, specificity, prevalence and, F1 score

```
accuracy_nb <- cm_nb$overall[["Accuracy"]]  
sensitivity_nb <- cm_nb$byClass[["Sensitivity"]]  
specificity_nb <- cm_nb$byClass[["Specificity"]]  
prevalence_nb <- cm_nb$byClass[["Prevalence"]]  
f1_nb <- cm_nb$byClass[["F1"]]
```

#Find the ROC and plot it. Show the AUC as well

```
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_nb) == "TRUE",1,0),m  
ethod="bin")  
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)  
plot(ciROC_bin95, col = 1, values=TRUE)  
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)
```



```
ROCit::ciAUC(pROC_bin)
```


```
##
## estimated AUC : 0.801438459045964
## AUC estimation method : binormal
##
## CI of AUC
## confidence level = 95%
## lower = 0.783049833697349 upper = 0.819827084394579
```

```
# translate income factor into binary outcome
temp <- adultpayclean_train %>%
  mutate(y = as.factor(income == "Above50K"))

#k-nearest neighbors with a train control and tuning
set.seed(2008)
# train control to use 10% of the observations each to speed up computations
control <- trainControl(method = "cv", number = 10, p = .9)
# train the model using knn. choose the best k value using tuning algorithm
train_knn <-
  train(
    y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship + education,
    method = "knn",
    data = temp,
    tuneGrid = data.frame(k = seq(3, 71, 2)),
    trControl = control
  )
```

```
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :
## non-uniform 'Rounding' sampler used
```

```
#plot the resulting model
ggplot(train_knn, highlight = TRUE)
```



```
#verify which k value was used  
train_knn$bestTune
```

```
##      k  
## 16 33
```

```
train_knn$finalModel
```

```
## 33-nearest neighbor model
## Training set outcome distribution:
##
## FALSE TRUE
## 19799 6453
```

```
#use this trained model to predict raw knn predictions
y_hat_knn <-
  predict(train_knn, adu1tpayclean_validation, type = "raw")

# compare the predicted and observed values using confusionMatrix to get the accuracy and other metrics
cm_knn <-
  confusionMatrix(y_hat_knn,
    as.factor(adu1tpayclean_validation$income == "Above50K"))
accuracy_knn <-
  confusionMatrix(y_hat_knn,
    as.factor(adu1tpayclean_validation$income == "Above50K"))$overall[["Accuracy"]]

cm_knn
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE 1992 360
##      TRUE   208 358
##
##           Accuracy : 0.8053
##           95% CI : (0.7905, 0.8196)
##      No Information Rate : 0.7539
##      P-Value [Acc > NIR] : 2.195e-11
##
##           Kappa : 0.4351
##
##      McNemar's Test P-Value : 2.361e-10
##
##           Sensitivity : 0.9055
##           Specificity : 0.4986
##           Pos Pred Value : 0.8469
##           Neg Pred Value : 0.6325
##           Prevalence : 0.7539
##           Detection Rate : 0.6827
##           Detection Prevalence : 0.8060
##           Balanced Accuracy : 0.7020
##
##           'Positive' Class : FALSE
##
```

```
#record the sensitivity, specificity, and prevalence
```

```
sensitivity_knn <- cm_knn$byClass[["Sensitivity"]]
```

```
specificity_knn <- cm_knn$byClass[["Specificity"]]
```

```
prevalence_knn <- cm_knn$byClass[["Prevalence"]]
```

```
f1_knn <- cm_knn$byClass[["F1"]]
```


```
#Find the ROC and plot it. Show the AUC as well
```

```
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_knn) == "TRUE",1,0),  
method="bin")
```

```
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)
```

```
plot(ciROC_bin95, col = 1, values=TRUE)
```

```
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)
```




```
ROCit::ciAUC(pROC_bin)
```

```
##  
## estimated AUC : 0.787212472866943  
## AUC estimation method : binormal  
##  
## CI of AUC  
## confidence level = 95%  
## lower = 0.763782743778267 upper = 0.810642201955618
```

```

#k-nearest classification using tuning function
set.seed(2008)

#train the model using knn3 classification
ks <- seq(3, 251, 2)
knntune <- map_df(ks, function(k) {
  temp <- adultpayclean_train %>%
    mutate(y = as.factor(income == "Above50K"))
  temp_test <- adultpayclean_validation %>%
    mutate(y = as.factor(income == "Above50K"))
  #create the knn3 model
  knn_fit <-
    knn3(
      y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship+education,
      data = temp,
      k = k
    )
  #predict the model for the current k
  y_hat <- predict(knn_fit, temp, type = "class")
  #get the confusionmatrix for the current k
  cm_train <- confusionMatrix(y_hat, temp$y)
  train_error <- cm_train$overall["Accuracy"]
  #do the same for test model
  y_hat <- predict(knn_fit, temp_test, type = "class")
  cm_test <- confusionMatrix(y_hat, temp_test$y)
  test_error <- cm_test$overall["Accuracy"]

  tibble(train = train_error, test = test_error)
})
#get the accuracy for the k with maximum accuracy
accuracy_knntune <- max(knntune$test)
#get the confusion matrix for that k
knn_fit <-
  knn3(
    y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship+education,
    data = temp,
    k = ks[which.max(knntune$test)]
  )
#predict the knn tune using the model for the k neighbor

```

```
y_hat_knntune <- predict(knn_fit, adultpayclean_validation, type = "class")
cm_knntune <- confusionMatrix(y_hat_knntune, as.factor(adultpayclean_validation$income == "Above50K"))

cm_knntune
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE 1994  351
##      TRUE   206  367
##
##           Accuracy : 0.8091
##           95% CI : (0.7944, 0.8232)
##      No Information Rate : 0.7539
##      P-Value [Acc > NIR] : 6.680e-13
##
##           Kappa : 0.448
##
##  Mcnemar's Test P-Value : 1.051e-09
##
##           Sensitivity : 0.9064
##           Specificity : 0.5111
##           Pos Pred Value : 0.8503
##           Neg Pred Value : 0.6405
##           Prevalence : 0.7539
##           Detection Rate : 0.6833
##           Detection Prevalence : 0.8036
##           Balanced Accuracy : 0.7088
##
##           'Positive' Class : FALSE
##
```

```

#record the sensitivity, specificity, and prevalence
sensitivity_knntune <- cm_knntune$byClass[["Sensitivity"]]
specificity_knntune <- cm_knntune$byClass[["Specificity"]]
prevalence_knntune <- cm_knntune$byClass[["Prevalence"]]
f1_knntune <- cm_knntune$byClass[["F1"]]

#Find the ROC and plot it. Show the AUC as well
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_knntune) == "TRUE",1
,0),method="bin")
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)
plot(ciROC_bin95, col = 1, values=TRUE)
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)

```



```
ROCit::ciAUC(pROC_bin)
```

```
##  
## estimated AUC : 0.794126916519226  
## AUC estimation method : binormal  
##  
## CI of AUC  
## confidence level = 95%  
## lower = 0.771082499427732 upper = 0.81717133361072
```

```
#recursive partitioning using rpart  
set.seed(2008)  
#train the model with the recursive partitioning  
train_rpart <-  
  train(  
    y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship+education,  
    method = "rpart",  
    tuneGrid = data.frame(cp = seq(0.0, 0.1, len = 25)),  
    data = temp  
  )
```

```
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used
```

```
#predict the outcomes with this model  
y_hat_rpart <- predict(train_rpart, adu1tpayclean_validation)  
#confusion matrix for the rpart model  
cm_rpart <-  
  confusionMatrix(y_hat_rpart,  
                  as.factor(adu1tpayclean_validation$income == "Above50K"))  
#get the accuracy  
accuracy_rpart <-  
  confusionMatrix(y_hat_rpart,  
                  as.factor(adu1tpayclean_validation$income == "Above50K"))$overall["Accuracy"]  
  
cm_rpart
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE  2002  324
##      TRUE   198  394
##
##           Accuracy : 0.8211
##           95% CI : (0.8067, 0.8349)
##      No Information Rate : 0.7539
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.4876
##
##  McNemar's Test P-Value : 4.472e-08
##
##           Sensitivity : 0.9100
##           Specificity : 0.5487
##      Pos Pred Value : 0.8607
##      Neg Pred Value : 0.6655
##           Prevalence : 0.7539
##      Detection Rate : 0.6861
##      Detection Prevalence : 0.7971
##      Balanced Accuracy : 0.7294
##
##      'Positive' Class : FALSE
##
```

```

#record the sensitivity, specificity, and prevalence
sensitivity_rpart <- cm_rpart$byClass[["Sensitivity"]]
specificity_rpart <- cm_rpart$byClass[["Specificity"]]
prevalence_rpart <- cm_rpart$byClass[["Prevalence"]]
f1_rpart <- cm_rpart$byClass[["F1"]]

#Find the ROC and plot it. Show the AUC as well
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_rpart) == "TRUE",1,0),method="bin")
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)
plot(ciROC_bin95, col = 1, values=TRUE)
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)

```



```
ROCit::ciAUC(pROC_bin)
```

```
##  
## estimated AUC : 0.815732110644685  
## AUC estimation method : binormal  
##  
## CI of AUC  
## confidence level = 95%  
## lower = 0.793892076612227 upper = 0.837572144677144
```

```
#random forest  
set.seed(2008)  
#train the vanilla random forest model  
train_rf <-  
  randomForest(y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship+education,  
               data = temp)  
  
y_hat_rf <- predict(train_rf, adu1tpayclean_validation)  
  
#create the confusionMatrix  
cm_rf <-  
  confusionMatrix(  
    y_hat_rf,  
    as.factor(adu1tpayclean_validation$income == "Above50K")  
  )  
#get the accuracy  
accuracy_rf <-  
  confusionMatrix(  
    y_hat_rf,  
    as.factor(adu1tpayclean_validation$income == "Above50K")  
  )$overall["Accuracy"]  
  
cm_rf
```




```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE  2012  332
##      TRUE   188  386
##
##           Accuracy : 0.8218
##           95% CI : (0.8074, 0.8355)
##      No Information Rate : 0.7539
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.4849
##
##      McNemar's Test P-Value : 3.588e-10
##
##           Sensitivity : 0.9145
##           Specificity : 0.5376
##           Pos Pred Value : 0.8584
##           Neg Pred Value : 0.6725
##           Prevalence : 0.7539
##           Detection Rate : 0.6895
##           Detection Prevalence : 0.8033
##           Balanced Accuracy : 0.7261
##
##           'Positive' Class : FALSE
##
```

```

#record the sensitivity, specificity, and prevalence
sensitivity_rf <- cm_rf$byClass[["Sensitivity"]]
specificity_rf <- cm_rf$byClass[["Specificity"]]
prevalence_rf <- cm_rf$byClass[["Prevalence"]]
f1_rf <- cm_rf$byClass[["F1"]]

#Find the ROC and plot it. Show the AUC as well
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_rf) == "TRUE",1,0),m
ethod="bin")
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)
plot(ciROC_bin95, col = 1, values=TRUE)
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)

```

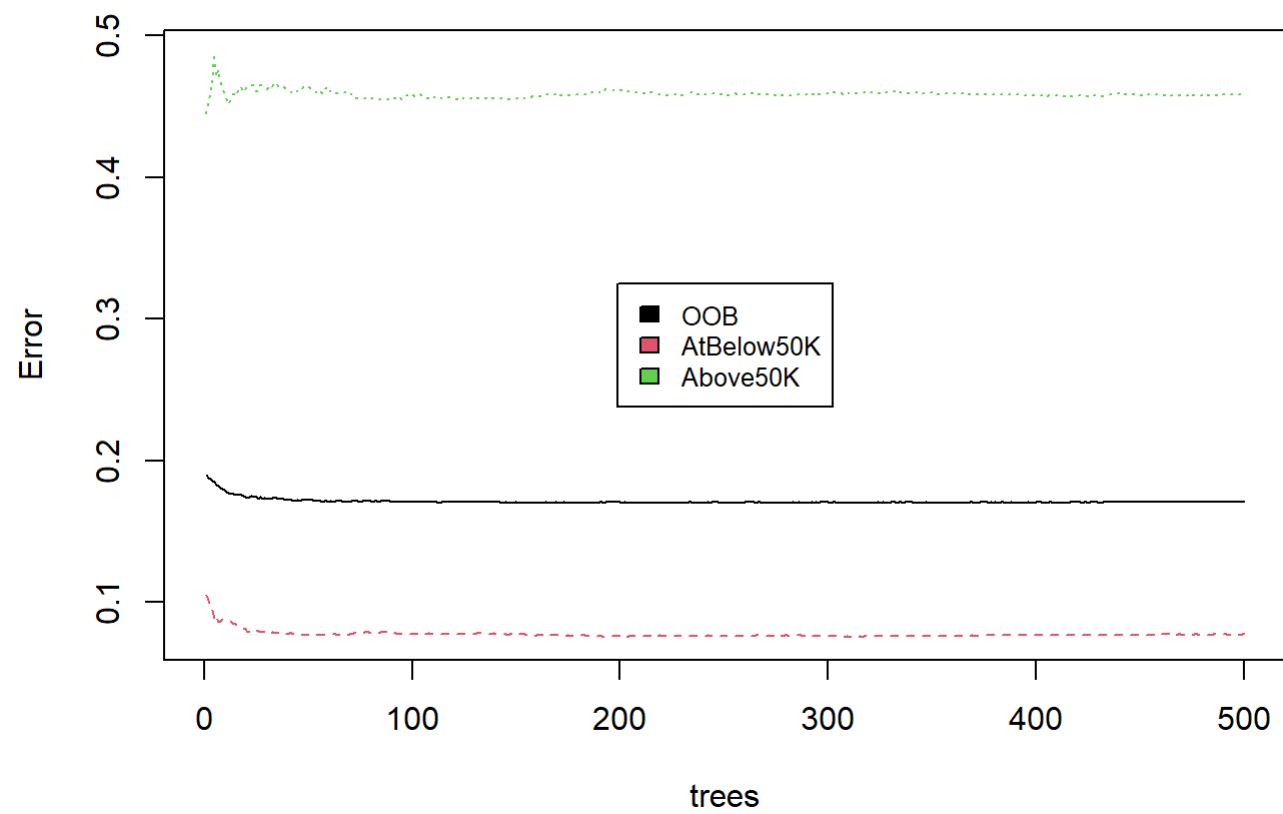


```
ROCit::ciAUC(pROC_bin)
```

```
##  
## estimated AUC : 0.818043983830059  
## AUC estimation method : binormal  
##  
## CI of AUC  
## confidence level = 95%  
## lower = 0.795999705923417 upper = 0.840088261736701
```

```
# Plot the error rate chart for the random forest  
plot(train_rf)  
legend("center", ifelse (colnames(train_rf$err.rate) == "FALSE", "AtBelow50K", ifelse (colnames(train_rf$err.rate) == "TRUE",  
"Above50K", "OOB")), col=1:4, cex=0.8, fill=1:4)
```

train_rf




```
set.seed(2008)
#random forest with tuning
nodesize <- seq(1, 90, 10)
acc <- sapply(nodesize, function(ns) {
  #train the model with tuning
  train(
    y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship + education + occupation + class,
    method = "rf",
    data = temp,
    tuneGrid = data.frame(mtry = 2),
    nodesize = ns
  )$results$Accuracy
})
```

```
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :
## non-uniform 'Rounding' sampler used
```

```
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used  
  
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used  
  
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used  
  
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used  
  
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used  
  
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used  
  
## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind = NULL) :  
## non-uniform 'Rounding' sampler used
```

```
qplot(nodesize, acc)
```



```
set.seed(2008)
#get the trained model for the max node size
train_rf_2 <-
  randomForest(
    y ~ age + eduyears + sex + race + hoursperweek + maritalstatus + relationship + education + occupation + class,
    data = temp,
    nodesize = nodesize[which.max(acc)]
  )
#predict the outcomes
y_hat_rf2 <- predict(train_rf_2, adu1tpayclean_validation)
#get the confusion matrix for random forest model
cm_rf2 <-
  confusionMatrix(
    y_hat_rf2,
    as.factor(adu1tpayclean_validation$income == "Above50K")
  )
#get the accuracy
accuracy_rftune <-
  confusionMatrix(
    y_hat_rf2,
    as.factor(adu1tpayclean_validation$income == "Above50K")
  )$overall["Accuracy"]

cm_rf2
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE  2024  291
##      TRUE   176  427
##
##           Accuracy : 0.84
##           95% CI : (0.8261, 0.8531)
##      No Information Rate : 0.7539
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.5441
##
##      McNemar's Test P-Value : 1.325e-07
##
##           Sensitivity : 0.9200
##           Specificity : 0.5947
##           Pos Pred Value : 0.8743
##           Neg Pred Value : 0.7081
##           Prevalence : 0.7539
##           Detection Rate : 0.6936
##           Detection Prevalence : 0.7934
##           Balanced Accuracy : 0.7574
##
##           'Positive' Class : FALSE
##
```

```
#record the sensitivity, specificity, and prevalence
```

```
sensitivity_rf2 <- cm_rf2$byClass[["Sensitivity"]]
```

```
specificity_rf2 <- cm_rf2$byClass[["Specificity"]]
```

```
prevalence_rf2 <- cm_rf2$byClass[["Prevalence"]]
```

```
f1_rf2 <- cm_rf2$byClass[["F1"]]
```


```
#Find the ROC and plot it. Show the AUC as well
```

```
pROC_bin <- ROCit::rocit(ifelse(adultpayclean_validation$income == "Above50K",1,0), ifelse(unname(y_hat_rf2) == "TRUE",1,0),  
method="bin")
```

```
ciROC_bin95 <- ROCit::ciROC(pROC_bin,level = 0.95)
```

```
plot(ciROC_bin95, col = 1, values=TRUE)
```

```
lines(ciROC_bin95$TPR~ciROC_bin95$FPR, col = 2, lwd = 2)
```

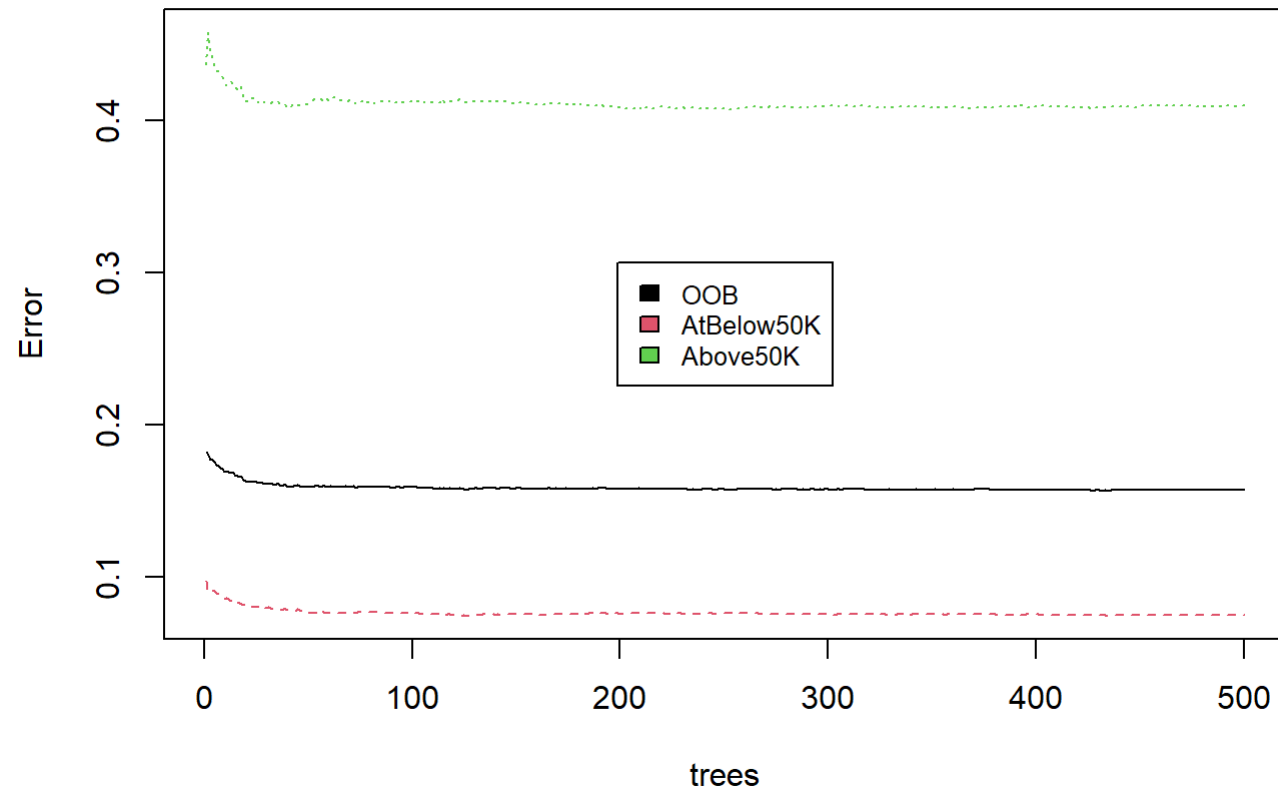


```
ROCit::ciAUC(pROC_bin)
```

```
##  
## estimated AUC : 0.849695675408264  
## AUC estimation method : binormal  
##  
## CI of AUC  
## confidence level = 95%  
## lower = 0.829656180065837 upper = 0.869735170750691
```

```
# Plot the error rate chart for the random forest  
plot(train_rf_2)  
legend("center", ifelse (colnames(train_rf_2$err.rate) == "FALSE", "AtBelow50K", ifelse (colnames(train_rf_2$err.rate) == "TRUE", "Above50K", "OOB")), col=1:4, cex=0.8, fill=1:4)
```

train_rf_2



tabulate all the accuracy results with sensitivity and specificity

```
accuracy_results <-  
  matrix(  
    c(  
      "Plain old guess",  
      round(accuracy_guess, 5),  
      round(sensitivity_guess, 5),  
      round(specificity_guess, 5),  
      round(prevalence_guess, 5),  
      round(f1_guess, 5),  
      "linear model",  
      round(accuracy_lm, 5),  
      round(sensitivity_lm, 5),  
      round(specificity_lm, 5),  
      round(prevalence_lm, 5),  
      round(f1_lm, 5),  
      "General linear model",  
      round(accuracy_glm, 5),  
      round(sensitivity_glm, 5),  
      round(specificity_glm, 5),  
      round(prevalence_glm, 5),  
      round(f1_glm, 5),  
      "naïve bayes",  
      round(accuracy_nb, 5),  
      round(sensitivity_nb, 5),  
      round(specificity_nb, 5),  
      round(prevalence_nb, 5),  
      round(f1_nb, 5),  
      "knn",  
      round(accuracy_knn, 5),  
      round(sensitivity_knn, 5),  
      round(specificity_knn, 5),  
      round(prevalence_knn, 5),  
      round(f1_knn, 5),  
      "knn tune",  
      round(accuracy_knntune, 5),  
      round(sensitivity_knntune, 5),  
      round(specificity_knntune, 5),  
      round(prevalence_knntune, 5),
```

```

round(f1_knntune, 5),
"rpart",
round(accuracy_rpart, 5),
round(sensitivity_rpart, 5),
round(specificity_rpart, 5),
round(prevalence_rpart, 5),
round(f1_rpart, 5),
"rf",
round(accuracy_rf, 5),
round(sensitivity_rf, 5),
round(specificity_rf, 5),
round(prevalence_rf, 5),
round(f1_rf, 5),
"rf tune",
round(accuracy_rftune, 5),
round(sensitivity_rf2, 5),
round(specificity_rf2, 5),
round(prevalence_rf2, 5),
round(f1_rf2, 5)
),
nrow = 9,
ncol = 6,
byrow = TRUE,
dimnames = list(
  c("1.", "2.", "3.", "4.", "5.", "6.", "7.", "8.", "9."),
  c(
    "Method",
    "Accuracy",
    "Sensitivity",
    "Specificity",
    "Prevalence",
    "F1"
  )
)
)
)
#style the table with knitr
accuracy_results %>% knitr::kable() %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed"))

```

	Method	Accuracy	Sensitlvity	Specifclity	Prevalence	F1
1.	Plain old guess	0.50822	0.50279	0.51	0.24606	0.33472
2.	linear model	0.81528	0.4429	0.93682	0.24606	0.54128
3.	General linear model	0.81049	0.38719	0.94864	0.24606	0.50135
4.	naive bayes	0.80089	0.91112	0.57218	0.67478	0.86064
5.	knn	0.80535	0.90545	0.49861	0.75394	0.87522
6.	knn tune	0.8098	0.90636	0.51114	0.75394	0.87745
7.	rpart	0.82111	0.91	0.54875	0.75394	0.88467
8.	rf	0.8218	0.91455	0.5376	0.75394	0.88556
9.	rf tune	0.83996	0.92	0.59471	0.75394	0.89657