

DatasetProcessingCode.R

rharidas

2021-12-03

```
#####
# Create Census pay train set, and validation set
#####

# Note: this process could take a couple of minutes

if (!require(tidyverse))
  install.packages("tidyverse", repos = "http://cran.us.r-project.org")

## Loading required package: tidyverse

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.2      v dplyr  1.0.7
## v tidyr   1.1.3      v stringr 1.4.0
## v readr   1.4.0      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

if (!require(caret))
  install.packages("caret", repos = "http://cran.us.r-project.org")

## Loading required package: caret

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

if (!require(data.table))
  install.packages("data.table", repos = "http://cran.us.r-project.org")
```

```

## Loading required package: data.table

##
## Attaching package: 'data.table'

## The following objects are masked from 'package:dplyr':
##
##     between, first, last

## The following object is masked from 'package:purrr':
##
##     transpose

if (!require(dplyr))
  install.packages("dplyr", repos = "http://cran.us.r-project.org")

if (!require(gridExtra))
  install.packages("gridExtra", repos = "http://cran.us.r-project.org")

## Loading required package: gridExtra

##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##     combine

if (!require(kableExtra))
  install.packages("kableExtra", repos = "http://cran.us.r-project.org")

## Loading required package: kableExtra

##
## Attaching package: 'kableExtra'

## The following object is masked from 'package:dplyr':
##
##     group_rows

if (!require(epiDisplay))
  install.packages("epiDisplay")

## Loading required package: epiDisplay

## Loading required package: foreign

## Loading required package: survival

```

```
##
## Attaching package: 'survival'

## The following object is masked from 'package:caret':
##
##      cluster

## Loading required package: MASS

##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##      select

## Loading required package: nnet

##
## Attaching package: 'epiDisplay'

## The following object is masked from 'package:lattice':
##
##      dotplot

## The following object is masked from 'package:ggplot2':
##
##      alpha
```

```
library(tidyverse)
library(caret)
library(data.table)
library(dplyr)
library(gridExtra)
library(kableExtra)
library(epiDisplay)

# Adult Census Income
# https://www.kaggle.com/uciml/adult-census-income

#download the dataset from the staging github location
dl <- tempfile()
download.file("https://github.com/rajeshharidas/havardxwork2/raw/main/adult.csv.zip",
             dl)

#read all the data into R dataset
adultpay <-
  fread(
    text = gsub(",", "\t", readLines(unzip(dl, "adult.csv"))),
    col.names = c(
      "age",
```

```

    "workclass",
    "fnlwgt",
    "education",
    "education.num",
    "marital.status",
    "occupation",
    "relationship",
    "race",
    "sex",
    "capital.gain",
    "capital.loss",
    "hours.per.week",
    "native.country",
    "income"
  )
)

#function to get the mode for categorical column. This is used to impute missing values
getmode <- function(v){
  v=v[nchar(as.character(v))>0]
  uniqv <- unique(v)
  uniqv[which.max(tabulate(match(v, uniqv)))]
}

#Keep only USA data
#Remove '?' from the work class and rename it to class, and finally remove workclass
#Rename all columns with a '.' in it
#Remove capital gain and loss column
#remove non-alphanumeric character from column data
#rename the label for below and above 50K income
#impute ? values to the modes in categorical columns
adultpayclean <- adultpay %>% filter (native.country == 'United-States') %>%
  mutate (class = ifelse(workclass == '?', getmode(adultpay$workclass), str_replace_all(workclass, "-",
dplyr::select(-workclass, -capital.gain, -capital.loss) %>%
  rename(
    c(
      edueyears = education.num,
      maritalstatus = marital.status,
      hoursperweek = hours.per.week,
      native = native.country
    )
  ) %>%
  mutate (maritalstatus = ifelse(
    maritalstatus == '?',
    getmode(adultpay$maritalstatus),
    str_replace_all(maritalstatus, "-", "")
  )) %>%
  mutate (occupation = ifelse(
    occupation == '?',
    getmode(adultpay$occupation),
    str_replace_all(occupation, "-", "")
  )) %>%
  mutate (education = ifelse(education == '?', getmode(adultpay$education), str_replace_all(education,

```

```

mutate (relationship = ifelse(
  relationship == '?',
  getmode(adultpay$relationship),
  str_replace_all(relationship, "-", ""))
)) %>%
mutate (native = ifelse(native == '?', 'Unknown', str_replace_all(native, "-", ""))) %>%
mutate (income = ifelse(
  income == '?',
  getmode(adultpay$income),
  str_replace_all(income, "<=50K", "AtBelow50K")
)) %>%
mutate (income = ifelse(
  income == '?',
  'Unknown',
  str_replace_all(income, ">50K", "Above50K")
))

# R 4.0 or later:
#convert all the character labels to factors
adultpayclean <-
as.data.frame(adultpayclean) %>% mutate(
  education = as.factor(education),
  maritalstatus = as.factor(maritalstatus),
  occupation = as.factor(occupation),
  relationship = as.factor(relationship),
  race = as.factor(race),
  sex = as.factor(sex),
  class = as.factor(class),
  income = as.factor(income)
)

# Validation set will be 10% of adultpay data
set.seed(1, sample.kind = "Rounding") # if using R 3.5 or earlier, use `set.seed(1)`

## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler
## used

```

```

test_index <-
createDataPartition(
  y = adultpayclean$income,
  times = 1,
  p = 0.1,
  list = FALSE
)
adultpayclean_train <- adultpayclean[-test_index, ]
adultpayclean_validation <- adultpayclean[test_index, ]

glimpse(adultpay)

```

```

## Rows: 32,561
## Columns: 15

```

```
## $ age          <int> 90, 82, 66, 54, 41, 34, 38, 74, 68, 41, 45, 38, 52, 32, ~
## $ workclass    <chr> "?", "Private", "?", "Private", "Private", "Private", "~
## $ fnlwgt       <int> 77053, 132870, 186061, 140359, 264663, 216864, 150601, ~
## $ education    <chr> "HS-grad", "HS-grad", "Some-college", "7th-8th", "Some--
## $ education.num <int> 9, 9, 10, 4, 10, 9, 6, 16, 9, 10, 16, 15, 13, 14, 16, 1~
## $ marital.status <chr> "Widowed", "Widowed", "Widowed", "Divorced", "Separated~
## $ occupation   <chr> "?", "Exec-managerial", "?", "Machine-op-inspct", "Prof~
## $ relationship <chr> "Not-in-family", "Not-in-family", "Unmarried", "Unmarri~
## $ race         <chr> "White", "White", "Black", "White", "White", "White", "~
## $ sex          <chr> "Female", "Female", "Female", "Female", "Female", "Fema~
## $ capital.gain  <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ capital.loss  <int> 4356, 4356, 4356, 3900, 3900, 3770, 3770, 3683, 3683, 3~
## $ hours.per.week <int> 40, 18, 40, 40, 40, 45, 40, 20, 40, 60, 35, 45, 20, 55, ~
## $ native.country <chr> "United-States", "United-States", "United-States", "Uni~
## $ income       <chr> "<=50K", "<=50K", "<=50K", "<=50K", "<=50K", "<=50K", "~
```

```
glimpse(adultpayclean)
```

```
## Rows: 29,170
## Columns: 13
## $ age          <int> 90, 82, 66, 54, 41, 34, 38, 74, 68, 45, 38, 52, 32, 51, ~
## $ fnlwgt       <int> 77053, 132870, 186061, 140359, 264663, 216864, 150601, 8~
## $ education    <fct> HSgrad, HSgrad, Somecollege, 7th8th, Somecollege, HSgrad~
## $ eduyears     <int> 9, 9, 10, 4, 10, 9, 6, 16, 9, 16, 15, 13, 14, 16, 15, 7, ~
## $ maritalstatus <fct> Widowed, Widowed, Widowed, Divorced, Separated, Divorced~
## $ occupation   <fct> Prof-specialty, Execmanagerial, Prof-specialty, Machineo~
## $ relationship <fct> Notinfamily, Notinfamily, Unmarried, Unmarried, Ownchild~
## $ race         <fct> White, White, Black, White, White, White, White, White, ~
## $ sex          <fct> Female, Female, Female, Female, Female, Female, Male, Fe~
## $ hoursperweek <int> 40, 18, 40, 40, 40, 45, 40, 20, 40, 35, 45, 20, 55, 40, ~
## $ native       <chr> "UnitedStates", "UnitedStates", "UnitedStates", "UnitedS~
## $ income       <fct> AtBelow50K, AtBelow50K, AtBelow50K, AtBelow50K, AtBelow5~
## $ class        <fct> Private, Private, Private, Private, Private, Private, Pr~
```

```
dim(adultpayclean)
```

```
## [1] 29170    13
```

```
dim(adultpayclean_train)
```

```
## [1] 26252    13
```

```
dim(adultpayclean_validation)
```

```
## [1] 2918     13
```

```
summary(adultpayclean)
```

```
##      age          fnlwgt          education          eduyears
## Min.   :17.00   Min.    : 12285   HSgrad       :9702   Min.    : 1.00
```

```

## 1st Qu.:28.00 1st Qu.: 115895 Somecollege:6740 1st Qu.: 9.00
## Median :37.00 Median : 176730 Bachelors :4766 Median :10.00
## Mean :38.66 Mean : 187069 Masters :1527 Mean :10.17
## 3rd Qu.:48.00 3rd Qu.: 234139 Assocvoc :1289 3rd Qu.:12.00
## Max. :90.00 Max. :1484705 11th :1067 Max. :16.00
## (Other) :4079
## maritalstatus occupation relationship
## Divorced : 4162 Execmanagerial:3735 Husband :11861
## MarriedAFspouse : 23 Profspecialty :3693 Notinfamily : 7528
## Marriedcivspouse :13368 Craftrepair :3685 Otherrelative: 696
## Marriedspouseabsent: 253 Admclerical :3449 Ownchild : 4691
## Nevermarried : 9579 Sales :3364 Unmarried : 3033
## Separated : 883 Otherservice :2777 Wife : 1361
## Widowed : 902 (Other) :8467
## race sex hoursperweek native
## Amer-Indian-Eskimo: 296 Female: 9682 Min. : 1.00 Length:29170
## Asian-Pac-Islander: 292 Male :19488 1st Qu.:40.00 Class :character
## Black : 2832 Median :40.00 Mode :character
## Other : 129 Mean :40.45
## White :25621 3rd Qu.:45.00
## Max. :99.00
## income class
## Above50K : 7171 Private :21794
## AtBelow50K:21999 Selfempnotinc: 2313
## Localgov : 1956
## Stategov : 1210
## Selfempinc : 991
## Federalgov : 886
## (Other) : 20

```

```
summary(adultpayclean_train)
```

```

## age fnlwgt education eduyears
## Min. :17.00 Min. : 12285 HSgrad :8716 Min. : 1.00
## 1st Qu.:28.00 1st Qu.: 116052 Somecollege:6071 1st Qu.: 9.00
## Median :37.00 Median : 176904 Bachelors :4318 Median :10.00
## Mean :38.66 Mean : 187117 Masters :1366 Mean :10.17
## 3rd Qu.:48.00 3rd Qu.: 234099 Assocvoc :1164 3rd Qu.:12.00
## Max. :90.00 Max. :1484705 11th : 948 Max. :16.00
## (Other) :3669
## maritalstatus occupation relationship
## Divorced : 3757 Execmanagerial:3382 Husband :10674
## MarriedAFspouse : 21 Profspecialty :3318 Notinfamily : 6803
## Marriedcivspouse :12033 Craftrepair :3300 Otherrelative: 629
## Marriedspouseabsent: 229 Admclerical :3095 Ownchild : 4213
## Nevermarried : 8616 Sales :3035 Unmarried : 2707
## Separated : 792 Otherservice :2490 Wife : 1226
## Widowed : 804 (Other) :7632
## race sex hoursperweek native
## Amer-Indian-Eskimo: 261 Female: 8708 Min. : 1.00 Length:26252
## Asian-Pac-Islander: 265 Male :17544 1st Qu.:40.00 Class :character
## Black : 2537 Median :40.00 Mode :character
## Other : 119 Mean :40.47

```


Dataset	Number of Rows	Number of Columns
train	26252	13
validation	2918	13

```
## White :23070 3rd Qu.:45.00
## Max. :99.00
##
## income class
## Above50K : 6453 Private :19587
## AtBelow50K:19799 Selfempnotinc: 2087
## Localgov : 1777
## Stategov : 1081
## Selfempinc : 910
## Federalgov : 793
## (Other) : 17
```

```
tribble(
  ~"Dataset", ~"Number of Rows", ~"Number of Columns",
  #-- /-- /----
  "train", nrow(adultpayclean_train), ncol(adultpayclean_train),
  "validation", nrow(adultpayclean_validation), ncol(adultpayclean_validation)
) %>% knitr::kable() %>% kable_styling(bootstrap_options = c("striped", "hover", "condensed"))
```


```
tab1(adultpayclean$education, sort.group = "decreasing", cum.percent = TRUE)
```

Distribution of adultpayclean\$education

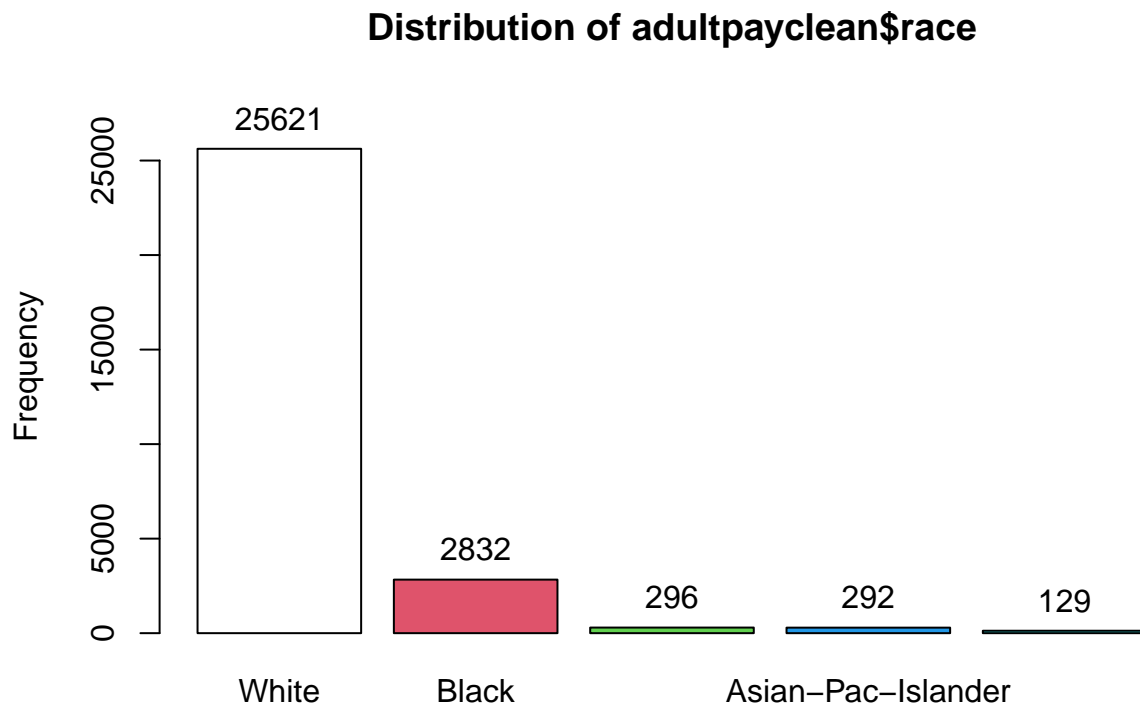



```
## adultpayclean$education :
##           Frequency Percent Cum. percent
## HSgrad          9702    33.3         33.3
## Somecollege      6740    23.1         56.4
## Bachelors        4766    16.3         72.7
## Masters          1527     5.2         77.9
## Assocvoc         1289     4.4         82.4
## 11th             1067     3.7         86.0
## Assocacdm         982     3.4         89.4
## 10th              848     2.9         92.3
## Profschool        502     1.7         94.0
## 7th8th            499     1.7         95.7
## 9th               395     1.4         97.1
## 12th              365     1.3         98.3
## Doctorate         328     1.1         99.5
## 5th6th            97      0.3         99.8
## 1st4th            46      0.2         99.9
## Preschool         17      0.1        100.0
## Total            29170   100.0        100.0
```

```
adultpayclean %>% group_by(education) %>%
  mutate(n=n()) %>% ggplot() +
  geom_bar(aes(education,col=income,fill=income)) + scale_y_sqrt()
```




```
tab1(adultpayclean$race, sort.group = "decreasing", cum.percent = TRUE)
```




```
## adultpayclean$race :
##               Frequency Percent Cum. percent
## White              25621      87.8         87.8
## Black               2832       9.7         97.5
## Amer-Indian-Eskimo   296       1.0         98.6
## Asian-Pac-Islander   292       1.0         99.6
## Other                129       0.4        100.0
## Total              29170     100.0        100.0
```

```
adultpayclean %>% group_by(race) %>%
  mutate(n=n()) %>% ggplot() +
  geom_bar(aes(race,col=income,fill=income)) + scale_y_sqrt()
```



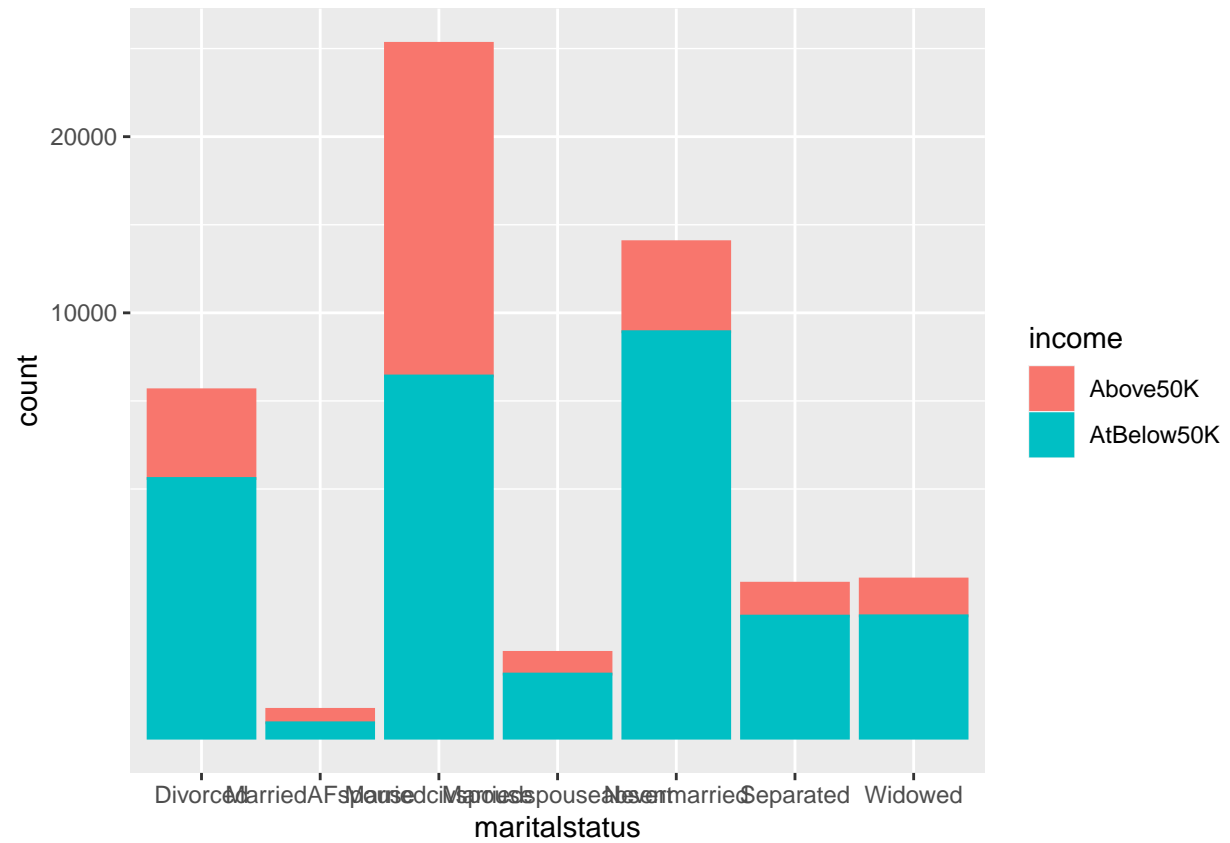
```
tab1(adultpayclean$maritalstatus, sort.group = "decreasing", cum.percent = TRUE)
```

Distribution of adultpayclean\$maritalstatus



```
## adultpayclean$maritalstatus :
##               Frequency Percent Cum. percent
## Marriedcivspouse      13368    45.8         45.8
## Nevermarried          9579    32.8         78.7
## Divorced              4162    14.3         92.9
## Widowed               902     3.1         96.0
## Separated             883     3.0         99.1
## Marriedspouseabsent    253     0.9         99.9
## MarriedAFspouse        23     0.1        100.0
## Total                29170   100.0        100.0
```

```
adultpayclean %>% group_by(maritalstatus) %>%
  mutate(n=n()) %>% ggplot() +
  geom_bar(aes(maritalstatus,col=income,fill=income)) + scale_y_sqrt()
```




```
tab1(adultpayclean$sex, sort.group = "decreasing", cum.percent = TRUE)
```

Distribution of adultpayclean\$sex




```
## adultpayclean$sex :  
##      Frequency Percent Cum. percent  
## Male      19488     66.8         66.8  
## Female     9682     33.2        100.0  
##   Total     29170    100.0        100.0
```

```
adultpayclean %>% group_by(sex) %>%  
  mutate(n=n()) %>% ggplot() +  
  geom_bar(aes(sex,col=income,fill=income)) + scale_y_sqrt()
```




```
tab1(adultpayclean$relationship, sort.group = "decreasing", cum.percent = TRUE)
```

Distribution of adultpayclean\$relationship



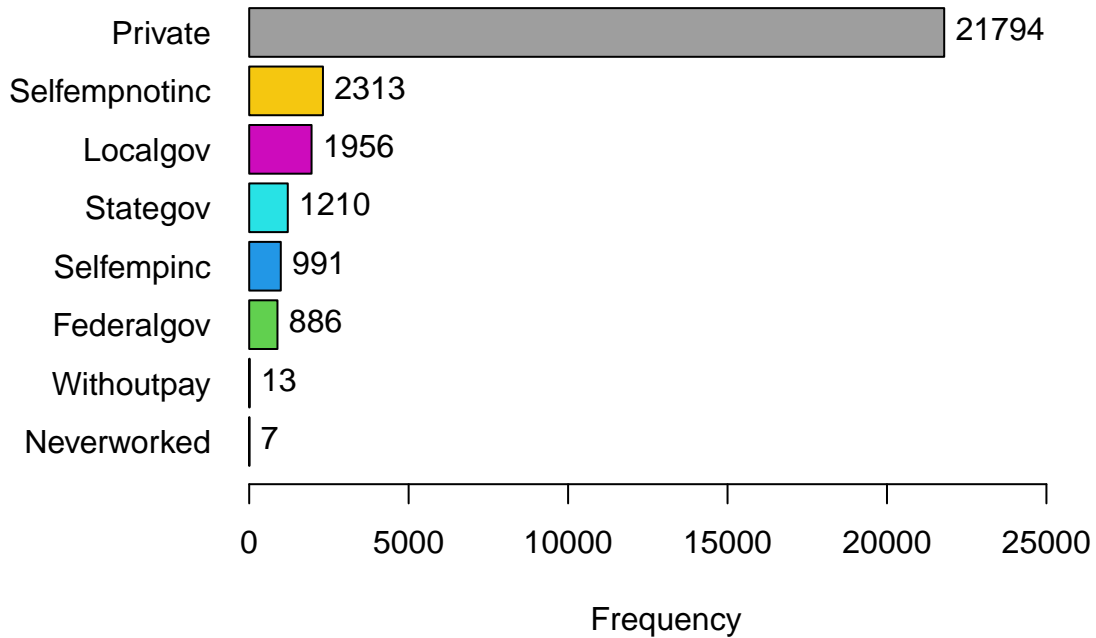
```
## adultpayclean$relationship :
##           Frequency Percent Cum. percent
## Husband           11861     40.7         40.7
## Notinfamily         7528     25.8         66.5
## Ownchild            4691     16.1         82.6
## Unmarried           3033     10.4         92.9
## Wife                1361      4.7         97.6
## Otherrelative         696      2.4        100.0
## Total              29170    100.0        100.0
```

```
adultpayclean %>% group_by(relationship) %>%
  mutate(n=n()) %>% ggplot() +
  geom_bar(aes(relationship,col=income,fill=income)) + scale_y_sqrt()
```


```
tab1(adultpayclean$class, sort.group = "decreasing", cum.percent = TRUE)
```

Distribution of adultpayclean\$class




```
## adultpayclean$class :
##           Frequency Percent Cum. percent
## Private           21794      74.7      74.7
## Selfempnotinc      2313       7.9      82.6
## Localgov           1956       6.7      89.3
## Stategov           1210       4.1      93.5
## Selfempinc          991       3.4      96.9
## Federalgov          886       3.0      99.9
## Withoutpay          13       0.0     100.0
## Neverworked         7       0.0     100.0
## Total             29170     100.0     100.0
```

```
adultpayclean %>% group_by(class) %>%
  mutate(n=n()) %>% ggplot() +
  geom_bar(aes(class,col=income,fill=income)) + scale_y_sqrt()
```



```
tab1(adultpayclean$income, sort.group = "decreasing", cum.percent = TRUE)
```

Distribution of adultpayclean\$income



```
## adultpayclean$income :  
##           Frequency Percent Cum. percent  
## AtBelow50K      21999      75.4         75.4  
## Above50K        7171      24.6        100.0  
## Total           29170     100.0        100.0
```

```
adultpayclean %>% group_by(age) %>%  
  mutate(n=n()) %>% ggplot() +  
  geom_bar(aes(age,col=income,fill=income)) + scale_y_sqrt()
```

